



# RITAL

Information retrieval and natural language processing  
Recherche d'information et traitement automatique de la langue

Master 1 DAC, semestre 2

Nicolas Thome



# Sequence Processing

**1** Sequence Processing

**2** Word Embeddings

- BoW and unsupervised methods (LSA/LDA): no explicit representation of words' sequences
- An important feature for many NLP task

⇒ How to include such sequential nature?

- Markov Models (MM) and HMM: generative models
- Conditional Random Fields (CRF): discriminative models → today
- RNNs & transformers: next course

## General framework

- Input  $\mathbf{x} \in \mathcal{X}$ : arbitrary
- Output  $y \in \mathcal{Y}$ : discrete output space
  - $\mathbf{y} \in \mathcal{Y}$  can be anything: vector, sequence (of vectors), graph, etc
- Assumption, there exist a function  $\psi(\mathbf{x}, \mathbf{y}) \in \mathbb{R}^d$
- CRF: model of conditional probability function:

$$P(\mathbf{y}|\mathbf{x}, \mathbf{w}) = \frac{e^{\mathbf{w}^T \psi(\mathbf{x}, \mathbf{y})}}{\sum_{\mathbf{y}' \in \mathcal{Y}} e^{\mathbf{w}^T \psi(\mathbf{x}, \mathbf{y}')}}}$$

- $Z(\mathbf{x}) = \sum_{\mathbf{y}' \in \mathcal{Y}} e^{\mathbf{w}^T \psi(\mathbf{x}, \mathbf{y}')}:$  partition function

## General framework for training

- A set of  $N$  training pairs  $\{(\mathbf{x}_i, \mathbf{y}_i)\}_{i \in \{1; N\}}$
- i.i.d samples, maximizing conditional log likelihood  $\mathcal{L}(\mathbf{w}) = \sum_{i=1}^N \log [P(\mathbf{y}_i | \mathbf{x}_i, \mathbf{w})]$

$$\mathcal{L}(\mathbf{w}) = \sum_{i=1}^N \mathbf{w}^T \psi(\mathbf{x}_i, \mathbf{y}_i) - \log \left( \sum_{\mathbf{y}' \in \mathcal{Y}} e^{\mathbf{w}^T \psi(\mathbf{x}_i, \mathbf{y}')} \right) = \mathbf{w}^T \psi(\mathbf{x}_i, \mathbf{y}_i) - \log (Z(\mathbf{x}_i)) \quad (1)$$

- $\mathbf{y}_i$  GT output for input  $\mathbf{x}_i$
- Eq (1) convex, can be solved with gradient descent
  - Can add  $\ell_2$  regularization  $\|\mathbf{w}\|^2$  on Eq (1) by putting Gaussian prior  $P(\mathbf{w}) \sim \mathcal{N}(0, \sigma^2)$

## Classification

- We need to define  $\mathcal{X}, \mathcal{Y}, \psi(\mathbf{x}, \mathbf{y})$
- Ex:  $\mathcal{X} = \mathbb{R}^d$ ,  $\mathcal{Y} = \{1; K\}$  ( $K$  number of classes), and  $\psi(\mathbf{x}, \mathbf{y}) \in \mathbb{R}^{d \cdot K}$

$$\psi(\mathbf{x}, \mathbf{y}) = (\delta_{y,1}\mathbf{x}, \dots, \delta_{y,k}\mathbf{x}, \dots, \delta_{y,K}\mathbf{x})$$

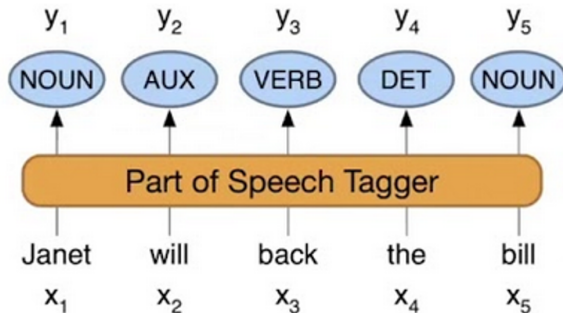
- $\mathbf{w} = (\mathbf{w}_1, \dots, \mathbf{w}_k, \dots, \mathbf{w}_K)$ , where  $\mathbf{w}_k \in \mathbb{R}^d \rightarrow \mathbf{w}^T \psi(\mathbf{x}, k) = \mathbf{w}_k^T \mathbf{x}$
- Training:  $\{(\mathbf{x}_i, k_i^*)\}_{i \in \{1; N\}}$

$$\Rightarrow P(k|\mathbf{x}, \mathbf{w}) = \frac{e^{\mathbf{w}_k^T \mathbf{x}}}{\sum_{k'=1}^K e^{\mathbf{w}_{k'}^T \mathbf{x}}}, \text{ and } \mathcal{L}(\mathbf{w}) = \sum_{i=1}^N \mathbf{w}_{k_i^*}^T \mathbf{x}_i - \log \left( \sum_{k=1}^K e^{\mathbf{w}_k^T \mathbf{x}_i} \right)$$

Logistic Regression

## Sequences: ex for Part of Speech tagging

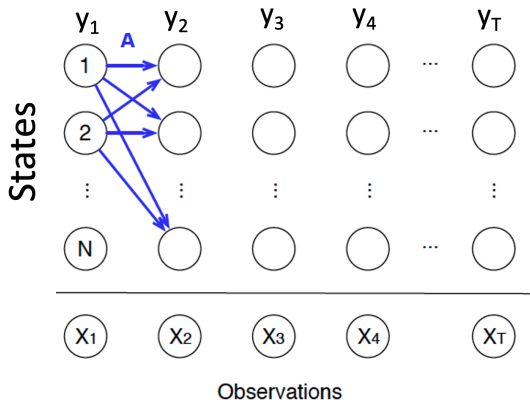
- Context: labelling of each word  $\rightarrow$  Part of Speech (PoS)
- Input  $\mathbf{x} = (x_1, \dots, x_T)$  a sentence with  $T$  words
- Output:  $\mathbf{y} = (y_1, \dots, y_T)$  a sequence of  $T$  PoS





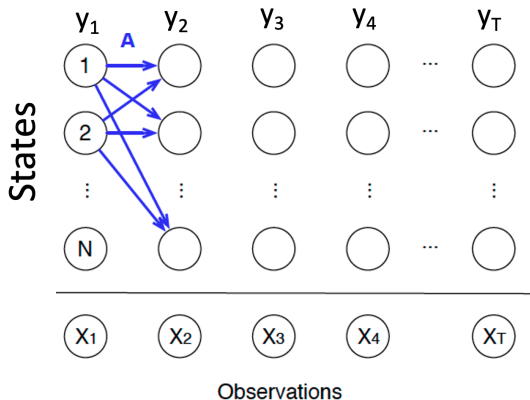
## HMMs

- Input  $\mathbf{x} = (x_1, \dots, x_T)$  a sentence with  $T$  words
  - $K$  words, e.g.  $K \sim 20000$
- Output:  $\mathbf{y} = (y_1, \dots, y_T)$  a sequence of  $T$  PoS
  - $N$  PoS tags, e.g.  $N = 20$
- HMM parameters: observation matrix  $B \in \mathbb{R}^{K \cdot N}$ , transition matrix  $A \in \mathbb{R}^{N \cdot N}$
- Training: states observed (not hidden)
  - A, B counting!
- Predicting PoS tags for a new sequence
  - Viterbi (see MAPSI)

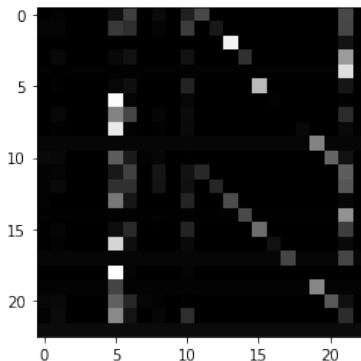


## HMMs

- Input  $\mathbf{x} = (x_1, \dots, x_T)$  a sentence with  $T$  words
  - $K$  words, e.g.  $K \sim 20000$
- Output:  $\mathbf{y} = (y_1, \dots, y_T)$  a sequence of  $T$  PoS
  - $N$  PoS tags, e.g.  $N = 20$
- HMM parameters: observation matrix  $B \in \mathbb{R}^{K \cdot N}$ , transition matrix  $A \in \mathbb{R}^{N \cdot N}$
- Training: states observed (not hidden)
  - A, B counting!
- Predicting PoS tags for a new sequence
  - Viterbi (see MAPSI)



- Some words are ambiguous: may be to several PoS tags
- Sequence order: help to resolve some ambiguities



## Examples of features for PoS Tagging

```
def features_full(sentence, index):
    return {
        'word': sentence[index],
        'is_first': index == 0,
        'is_last': index == len(sentence) - 1,
        'is_capitalized': sentence[index][0].upper() == sentence[index][0],
        'is_all_caps': sentence[index].upper() == sentence[index],
        'is_all_lower': sentence[index].lower() == sentence[index],
        'prefix-1': sentence[index][0],
        'prefix-2': sentence[index][:2],
        'prefix-3': sentence[index][:3],
        'suffix-1': sentence[index][-1],
        'suffix-2': sentence[index][-2:],
        'suffix-3': sentence[index][-3:],
        'prev_word': 'if index == 0 else sentence[index - 1],
        'next_word': 'if index == len(sentence) - 1 else sentence[index + 1],
        'has_hyphen': '-' in sentence[index],
        'is_numeric': sentence[index].isdigit(),
        'capitals_inside': sentence[index][1:].lower() != sentence[index][1:]
    }
```

$$P(\mathbf{y}|\mathbf{x}) = \frac{e^{\sum_{t=1}^T \sum_{k=1}^K [\theta_k u_k(y_t, \mathbf{x}) + \lambda_k p_k(y_{t-1}, y_t, \mathbf{x})]}}{Z(\mathbf{x})}$$
$$Z(\mathbf{x}) = \sum_{\mathbf{y}' \in \mathcal{Y}} e^{\sum_{t=1}^T \sum_{k=1}^K [\theta_k u_k(y'_t, \mathbf{x}) + \lambda_k p_k(y'_{t-1}, y'_t, \mathbf{x})]} \quad (2)$$

- Problem:  $\mathcal{Y}$  huge:  $\dim N^T$
- Model prediction:  $\hat{\mathbf{y}} = \arg \max_{\mathbf{y}} P(\mathbf{y}|\mathbf{x})$
- Training : computing normalization factor  $Z(\mathbf{x}) = \sum_{\mathbf{y}' \in \mathcal{Y}} \dots$
- Brute-force computation intractable, must exploit structure...

$\Rightarrow$  dynamic programming

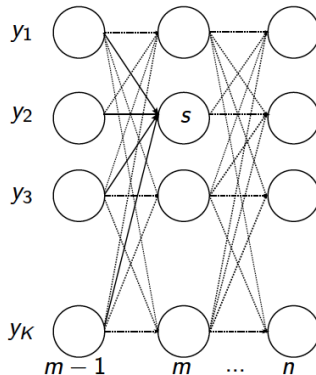
## Inference

- Model prediction:

$$\begin{aligned}
 \hat{\mathbf{y}} &= \arg \max_{\mathbf{y}} P(\mathbf{y}|\mathbf{x}) \\
 &= \arg \max_{\mathbf{y}} \sum_{t=1}^T \sum_{k=1}^K F_k(y_{t-1}, y_t, \mathbf{x}) \\
 &:= \arg \max_{\mathbf{y}} \sum_{t=1}^T g_t(y_{t-1}, y_t)
 \end{aligned}$$

- Viterbi algorithm  $\sim$  HMM

$$\delta_m(s) \triangleq \max_{\{y_1, \dots, y_{m-1}\}} \left[ \sum_{i=1}^{m-1} g_i(y_{i-1}, y_i) + g_m(y_{m-1}, s) \right]$$



## Training

$$\mathcal{L}(\mathbf{w}) = \mathbf{w}^T \psi(\mathbf{x}_i, \mathbf{y}_i) - \log(Z(\mathbf{x}_i))$$

- $Z(\mathbf{x}_i)$  intractable
- Forward-backward algorithm

---

**Algorithm 1** The forward-backward algorithm for the probability calculation of the CRF

---

**Input:** The model  $P(Y|X)$ , the input sequence  $\mathbf{x}$ , the output sequence  $\mathbf{y}$ , and the location  $i$

**Output:** The conditional probabilities  $P(Y_i = y_i | \mathbf{x})$ ,  $P(Y_{i-1} = y_{i-1}, Y_i = y_i | \mathbf{x})$

1. Let  $M_i(y_{i-1}, y_i | \mathbf{x}) = \exp(\sum_{k,k} \lambda_k t_k(y_{i-1}, y_i, \mathbf{x}, i) + \sum_{j,j} \mu_j s_j(y_i, \mathbf{x}, i))$ ,  $y_0 = start$ ,  $y_{n+1} = stop$

2. Initialization,  $\alpha_0(y_0 | \mathbf{x}) = 1$ ,  $\beta_{n+1}(y_{n+1} | \mathbf{x}) = 1$

3. Recursion, for  $k = 1, 2, \dots, i$

$$\alpha_k^T(y_k | \mathbf{x}) = \alpha_{k-1}^T(y_{k-1} | \mathbf{x}) M(y_{k-1}, y_k | \mathbf{x})$$

4. Recursion, for  $j = n, n-1, \dots, i+1, i, i-1, \dots, 1$

$$\beta_j(y_j | \mathbf{x}) = M_{j+1}(y_j, y_{j+1} | \mathbf{x}) \beta_{j+1}(y_{j+1} | \mathbf{x})$$

5. Calculation,  $Z(\mathbf{x}) = \mathbf{1}^T \times \beta_1(\mathbf{x})$

6. Calculation,  $P(Y_i = y_i | \mathbf{x}) = \alpha_i^T(y_i | \mathbf{x}) \beta_i(y_i | \mathbf{x}) / Z(\mathbf{x})$

$$P(Y_{i-1} = y_{i-1}, Y_i = y_i | \mathbf{x}) = \alpha_{i-1}^T(y_{i-1} | \mathbf{x}) M_i(y_{i-1}, y_i | \mathbf{x}) \beta_i(y_i | \mathbf{x}) / Z(\mathbf{x})$$


---

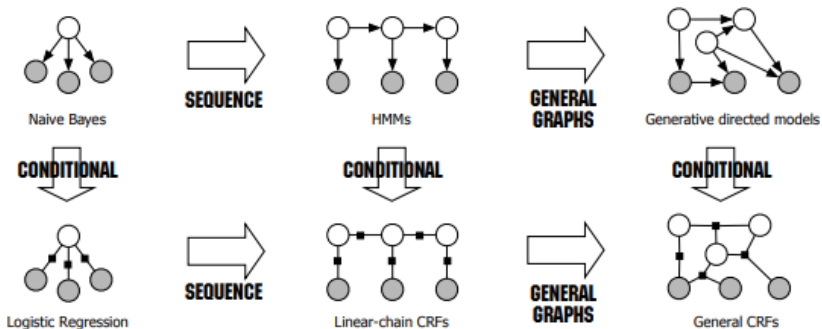


Fig. 2.3 Diagram of the relationship between naive Bayes, logistic regression, HMMs, linear-chain CRFs, generative models, and general CRFs.



# Word Embeddings

1 Sequence Processing

2 Word Embeddings